

Normalization

The Normal Form Hierarchy

- A top-down approach to relational design (as opposed to ER diagramming, which is bottom-up)
- 1NF, 2NF, 3NF, BCNF, 4NF, 5NF
- Higher normal forms are more restrictive.
- There is a greater number of tables as you go up to a higher normal form.
- If a table is in a higher NF, then it is in all of the lower NFs.
 - (Only by convention, really, as it is possible, e.g., for a relation to be consistent w/ 3NF def and not w/ 2NF)
- Note that:
 - We strive to achieve 3NF
 - If you use the ER model properly and use the translation techniques properly, the relational schema are usually in 3NF.

Update Anomalies

- Update anomalies are one of the problems we are looking to avoid by employing normalization.
- They result from uncontrolled redundancy resulting from poor design decisions.
- Note that both tables contain two separate concepts
- Updating a value of, for example, Dname, requires changing multiple records. If Dname is not changed for all employees in the department in question, an *update anomaly* occurs

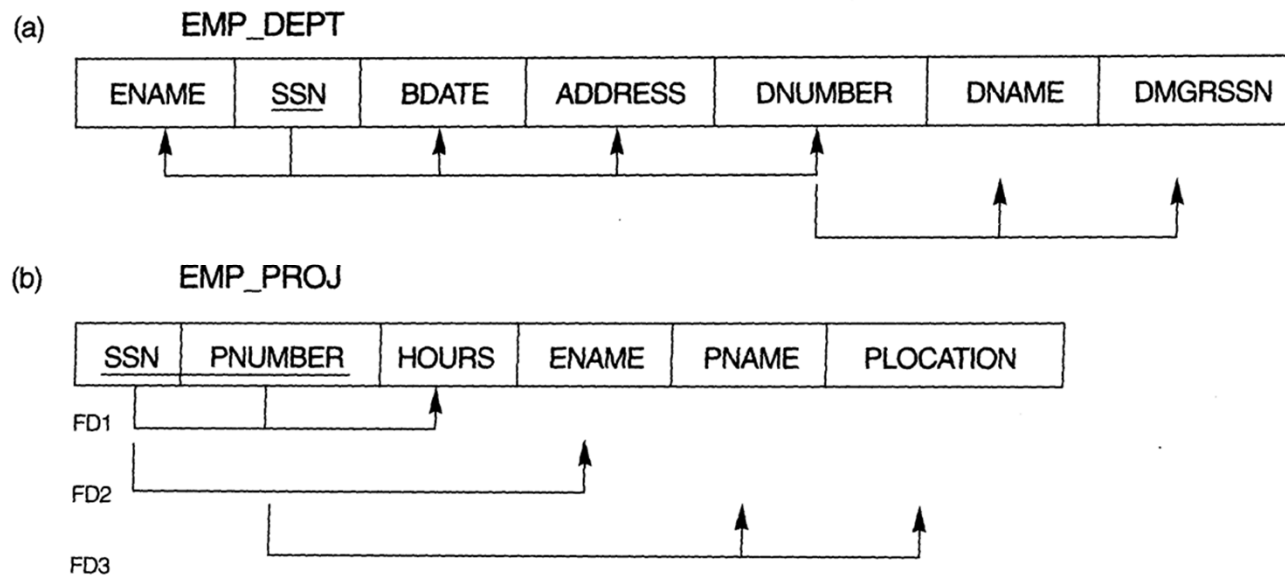


FIGURE 10.3 Two relation schemas suffering from update anomalies.

Update Anomaly Types

- Insert Anomalies
 - Can't insert a new department into emp_dept above that has no employees yet.
 - Every new employee requires consistent re-entry of all details for that employee's department.
- Deletion Anomalies
 - What happens when you delete the last employee from the R&D dept? All data about that department is gone(!)
- Modification Anomalies
 - If we change some info about a department, we have to change it consistently across all employees who work for that department
 - (Not as big a deal b/c a single SQL query will still do the job.)
- Related problem: null values proliferate in poorly designed relations (eg employees with no departments)

Spurious Tuples

- Tuples that result from a ‘lossy join’
 - Lossy join: joined table contains tuples that did not exist in underlying tables
- Decomposition per normalization avoids lossy joins.
- Example:
- Base Relation:
 - Emp-Proj (SSN, P#, Hours, Ename, Pname, Ploc)
 - (111-11-1111, P12, 40, Bill, ProjX, Houston)
 - (222-22-2222, P37, 20, Mary, ProjY, Houston)
- Suppose we decompose to:
 - R1(Ename, Ploc)
 - (Bill, Houston)
 - (Mary, Houston)
 - R2(SSN, P#, Hours, Pname, Ploc)
 - (111-11-1111, P12, 40, ProjX, Houston)
 - (222-22-2222, P37, 20, ProjY, Houston)
- What happens when we do a natural join?
 - We gain nonsense (‘spurious’) tuples
 - Data integrity is lost after join → ‘lossy’ join

Functional Dependencies

- $X \rightarrow Y$
- “X functionally determines Y”
- For a given value of X there can be one and only one value of Y
- If you know X, you also know with certainty Y
- If you know a person’s SSN, you can determine for sure their first name
- If you know a person’s first name you CANNOT determine their SSN (b/c hundreds of thousand of people, all w/different SSNs, may share that name!)

Functional Dependencies

- B is a function of A if for every A there is at most one value for B
 - Sound familiar?
 - Similar to cardinality, but attribute is unit of measure
 - Representation:
 - $A \rightarrow B$
 - Interpretation:
 - “B is a function of A” or
 - “A functionally determines B”

Two Entities, A & B

- Cardinality Analogy: Functional Dependencies
 - If A:B are 1:1
 $A \rightarrow B \ \& \ B \rightarrow A$
 - If A:B are 1:N
 $B \rightarrow A \text{ ONLY}$
 - If A:B are N:1
 $A \rightarrow B \text{ ONLY}$
 - If A:B are N:M
There are NO functional dependencies
- Note, though, that FD pertains to *attributes* where cardinality pertains to *relationships*

FD Inferences

- An FD is a property of the relational schema, NOT any particular state of the relation
 - Have to have a priori knowledge of the semantics to know which FDs hold
 - (Just like you did with cardinality, btw)
 - HOWEVER, FDs can be ruled *out* by looking at the data.

FD Example

- Which FDs can be ruled *out* here?

TEACH

TEACHER	COURSE	TEXT
Smith	Data Structures	Bartram
Smith	Data Management	Al-Nour
Hall	Compilers	Hoffman
Brown	Data Structures	Augenthaler

Normalization Terminology

- Superkey – set of attributes assuring uniqueness
 - (any set that contains one or more candidate keys)
- Candidate key – *minimal* set of attributes assuring uniqueness
- Primary Key – one set of attributes arbitrarily chosen from among set of candidate keys to act as PK
- Prime attribute – member of some candidate key
- Nonprime attribute – *not* a member of any candidate key

First Normal Form

- A relation R is in 1 NF if all attributes have atomic values

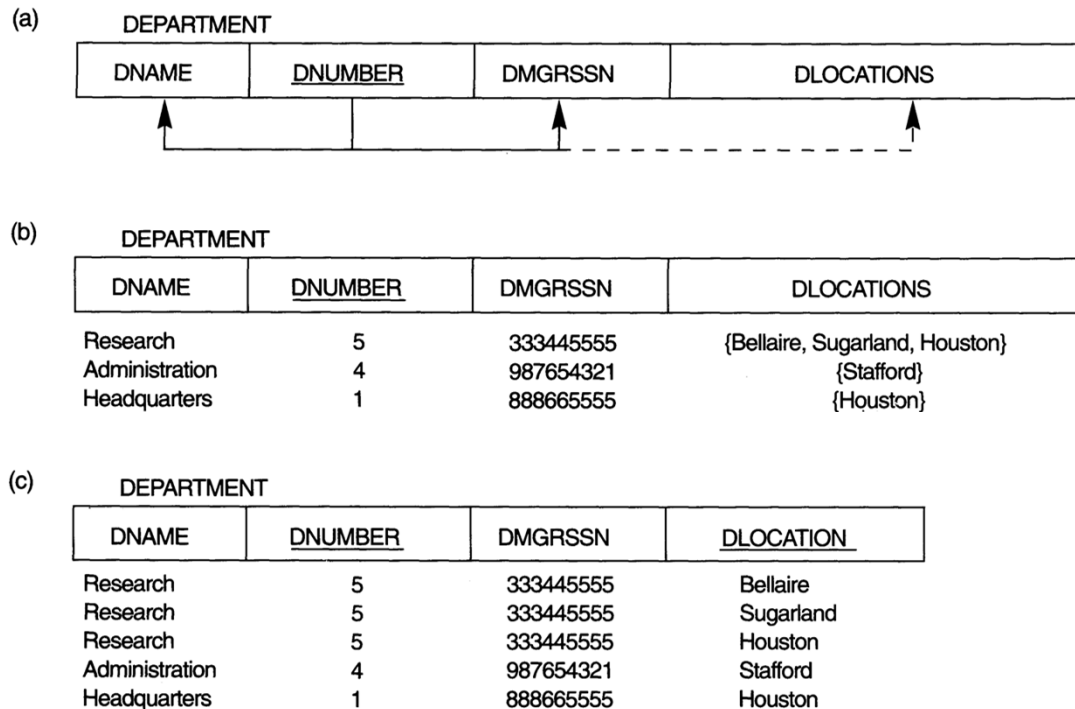


Figure 14.8 Normalization into 1NF. (a) Relation schema that is not in 1NF.

(b) Example relation instance. (c) 1NF relation with redundancy.

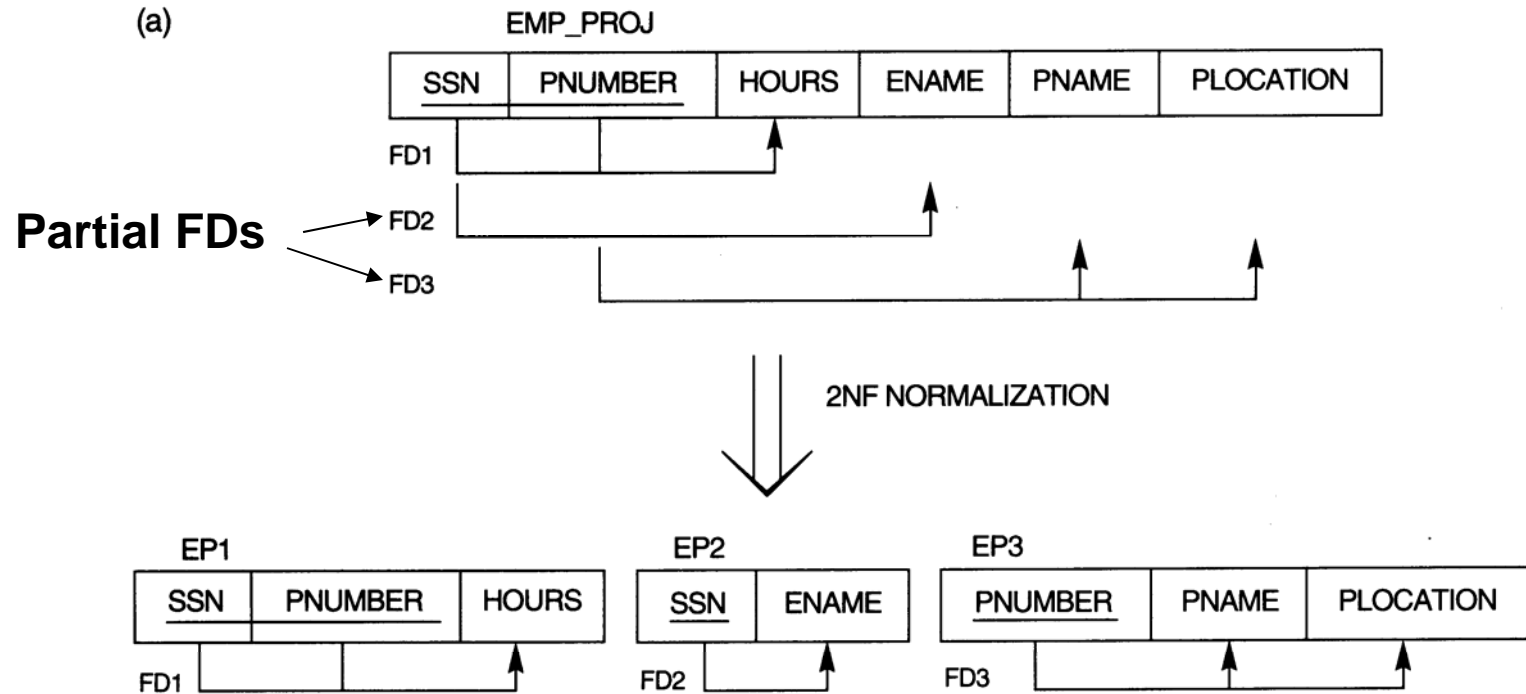
- Either accept redundancy, like c) above, or (far better) convert to a multi-valued attribute weak entity situation.
- (Fairly) recent trends – object DBs, XML allows nesting that violated 1NF
- B is *NOT* in 1 NF, C *IS* in 1 NF

Second Normal Form

- A relation R is in 2 NF if every nonprime attribute A in R is *fully* functionally dependent on the primary key of R.
 - I.e., 2NF relations have no *partial dependencies*
 - Issue only with relations w/ composite PKs
 - The test for partial dependencies:
 - Take away a part of the concatenated key, if any FD still holds, then there is a *partial dependency*
 - » *Relation is NOT in 2NF*

2NF Example

(a)



Third Normal Form

- Definition: A relation is in 3 NF if it satisfied 2NF conditions and it has no *transitive dependencies*:
- $X \rightarrow Z$, but only because $X \rightarrow Y$ and $Y \rightarrow Z$ is a *transitive dependency*
- EX:
- $Emp\#, EmpName, Phone, Salary, Department\#, Dname, Dlocation$
- Not in 3NF because $Emp\# \rightarrow Dname$ or $Dloc$ only because $Emp\# \rightarrow Dept\#$ and $Dept\#$ in turn $\rightarrow Dname, Dloc$

3NF Example

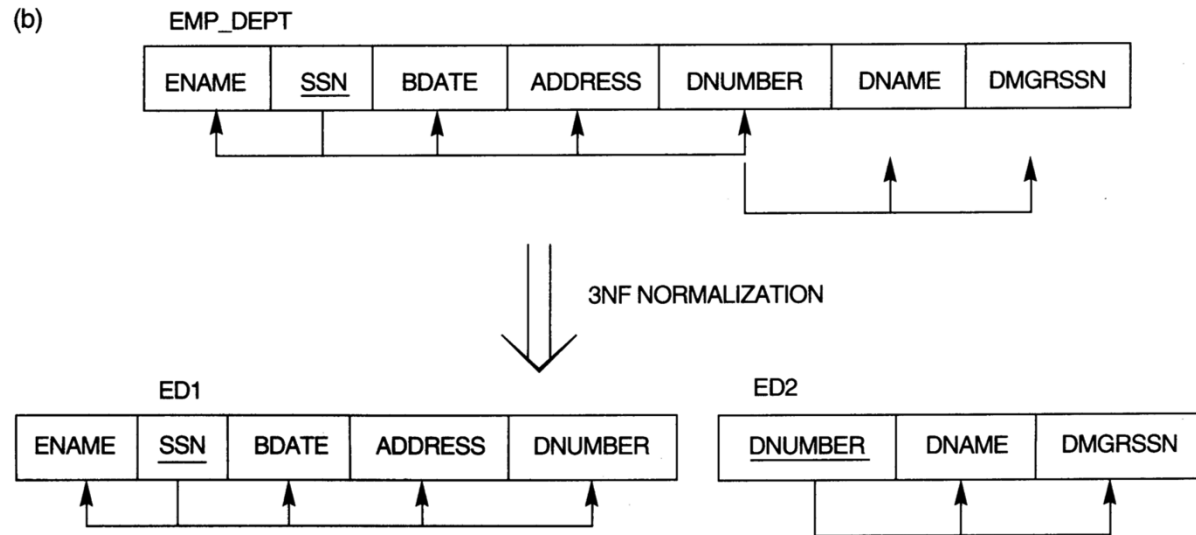


Figure 14.10 The normalization process. (a) Normalizing EMP_PROJ into 2NF relations. (b) Normalizing EMP_DEPT into 3NF relations.

- Dnumber → Dname
- Dnumber → Dmgrssn
- Dnumber is not a key
- Therefore (b) above is NOT in 3NF

3NF Issues

- If relation is not in 3nf, susceptible to update anomalies
- Note that a Non 3NF table typically contains two concepts that are in a 1:N relationship as in the EMP_Dept (N:1) example
- Remedy:
 - Create a separate relation for each set of directly dependent attributes
- If you employ ER modeling properly and translate correctly, the resulting schema will be in 3NF.
- In practice, we usually stop at 3NF as higher NFs do not occur frequently.

Normalizing to 3NF

- Every non-key attribute must provide a fact about “the key, the whole key, and nothing but the key”
- Give attributes with transitive dependencies their own table

The Bottom Line

- If you apply the rules of semantic modeling that we learned in class (i.e. our process for ER diagramming and Translation), your resultant relational schema should *be in 3rd normal form anyway!*
- Higher NF = More tables
- More tables =
 - Degraded performance
 - Query complexity
 - Referential integrity challenges
- Practical Approach
 - Use semantic approach like ER model
 - Convert to schema
 - Check for 3NF
 - If absent, normalize to 3NF
 - Ignore normalization wisely, consciously, typically in the name of performance: *denormalization*

Other Normal Forms

- Normal forms up to 6th, plus Boyce-Codd Normal Form exist, but are rare in practice and beyond the scope of this class
- See http://en.wikipedia.org/wiki/Database_normalization#Normal_forms

for a good discussion of these forms if interested.

(Optional) See the following slides for BCNF, 4NF, 5NF, very informally presented

Higher Normal Forms

- BCNF, 4NF, 5NF
- Hard to recognize, unimportant in practice
- BCNF:
 - If $X \rightarrow Y$ in R, then X must be a superkey of R
- Example:
 - Property (Prop#, County, Lot#, Area)
 - $\text{Prop\#} \rightarrow \text{County, Lot\#, Area}$
 - $\text{Area} \rightarrow \text{County}$
 - Area is not a superkey but it IS prime, so OK per 3NF, violate BCNF
 - (Prop#, Area, Lot#) & (Area, County)

4NF

- Involves multi-valued dependencies
- Emp (Ename, Pname, Dname)
 - Employee can have multiple dependents and work on multiple departments
 - Break out to
 - (Ename, Pname) and (Ename, Dname)

5 NF

- Pertains to ternary relationships
- Looks to avoid “join dependencies”
- Supply (Sup#, Part#, Proj#)
- IF it is the case that:
 - For a supplier S that supplies part P and
 - A project X that uses part P and
 - Supplier S supplies at least one part to project X
 - Then supplier S will also necessarily supply part P to project X
- THEN Supply is NOT in 5NF and should be
 - (Sup#, Part#), (Sup#, Proj#), (Part#, Proj#)
- Review from ERD perspective
- Obviously requires intuition/skill to identify JDs
 - Really, really not worried about in practice